

👉 XAI: Evaluating Interpretability

FAMAF - UNC

November 4, 2025



**eXplainable
Artificial Intelligence**

Disclaimer ---

This course is based on

Explainable Artificial Intelligence

From Simple Predictors to Complex Generative Models

Spring 2023, Harvard University

<https://interpretable-ml-class.github.io/>


Overview ---

- Evaluating interpretability in the interpretable ML community:
 - ▶ Interpretability depends on human experience of the model
 - ▶ Disagreement about the best way to measure it
- These papers:
 - ▶ Evaluating factors related to interpretability through user studies



🍷 Other Relevant Fields ---

- **Human-Computer Interaction (HCI):**
 - ▶ Theories for how people **interact with technology**
- **Psychology:**
 - ▶ Theories for how people **process information**
- Both have thought carefully about **experimental design**



Outline ---

- Research paper: **“Human Evaluation of Models Built for Interpretability”** by Lage et al.
- Research paper: **“Manipulating and Measuring Model Interpretability”** by Poursabzi-Sangdeh et al.
- Discussion

The Seventh AAAI Conference on Human
Computation and Crowdsourcing (HCOMP-19)

Human Evaluation of Models Built for Interpretability

Isaac Lage,^{*1} Emily Chen,^{*1} Jeffrey He,^{*1} Menaka Narayanan,^{*1}
Been Kim,² Samuel J. Gershman,¹ Finale Doshi-Velez¹

¹Harvard University, ²Google

isaaclage@g.harvard.edu, {emily-chen, jdhe, menakanarayanan}@college.harvard.edu
beenkim@google.com, gershman@fas.harvard.edu, finale@seas.harvard.edu

Abstract

Recent years have seen a boom in interest in interpretable machine learning systems built on models that can be understood, at least to some degree, by domain experts. However, exactly what kinds of models are truly human-interpretable remains poorly understood. This work advances our understanding of precisely which factors make models inter-

human-simulatable models include human-simulatable regression functions (e.g. Caruana et al., 2015; Ustun and Rudin, 2016), and various logic-based methods (e.g. Wang and Rudin, 2015; Lakkaraju, Bach, and Leskovec, 2016; Singh, Ribeiro, and Guestrin, 2016).

But even types of models designed to be human-simulatable may cease to be so when they are too complex



Definitions ---

- **Human-simulatability:** The ability of humans to mentally trace through a model's decision-making process and accurately predict its outputs for given inputs, essentially simulating how the model would behave.
- **Decision set:** A logic-based machine learning model consisting of a collection of independent logical rules, where each rule maps a combination of input features to an output prediction.
- **Decision set complexity:** A decision set is more complex when it has more rules, more conditions per rule, or conditions that are more cognitively difficult to evaluate.

Contributions ---

Research Questions:

- Which types of **decision set complexity** most affect human-simulatability?
- Is relationship between complexity and human-simulatability **context dependent**?

Approach:

- Large scale, carefully controlled user studies

Decision Sets ---

- **Logic-based models** are often considered interpretable
- Many approaches for **learning them from data** (Decision Trees)

faithful and cold or **brave and passive** → **candy or dairy** and **fruit**
sleepy or **patient and obedient** → **spices and grains** or **dairy**
brave and **sleepy or patient or laughing** → **dairy and fruit** or **grains**
crying or **sleepy and faithful** → **grains** and **spices or fruit**


🍷 Regularizers ---

- **Regularizer (paper context):** A technique that penalizes certain types of complexity during decision set training (such as number of rules, conditions per rule, or specific types of conditions) to make the models **less complex and more interpretable**.
- What kinds of complexity is it **most urgent to regularize** to learn interpretable models?



🍲 Types of Complexity ---

- **Model size:** Number of rules/lines in the decision set
- **Variable repetitions:** Same variable appearing multiple times
- **Cognitive chunks:** Number of distinct logical conditions



What if we optimized the models with data?

🍲 Context: Domains ---

Ingredients:

- **Vegetables:** okra, carrots, spinach
- **Spices:** turmeric, thyme, cinnamon
- **Dairy:** milk, butter, yogurt
- **Fruit:** mango, strawberry, guava
- **Candy:** chocolate, taffy, caramel
- **Grains:** bagel, rice, pasta



Low Risk: Alien meal recommendation

Disease Medications:

- **antibiotics:** Aerove, Adenon, Athoxin
- **painkillers:** Poxin, Parola, Pelapin
- **vitamins:** Vipryl, Vyorix, Votasol
- **stimulants:** Silvax, Setoxin, Soderal
- **tranquilizers:** Trasmin, Tydesol, Texopal
- **laxatives:** Lantone, Lezanto, Lexerol



High Risk: Alien medical prescription

What if we used 2 different real domains?

Context: Tasks ---

bubbly or clumsy → brave

faithful and cold or brave and passive → candy or dairy and fruit

sleepy or patient and obedient → spices and grains or dairy

brave and sleepy or patient or laughing → dairy and fruit or grains

crying or sleepy and faithful → grains and spices or fruit

Observations: patient, wearing glasses, lazy

Recommendation: milk, guava

- **Simulation:**

- ▶ What would the model recommend the alien?

- **Verification:**

- ▶ Is **milk and guava** a correct recommendation?

- **Counterfactual:**

- ▶ If **patient** were replaced with **sleepy**, would the correctness of the **milk and guava** recommendation change?

What if we used more realistic tasks?

🍷 Tradeoff Between Control and Generalizability ---

Tradeoff between the ability to **tightly control** the experiment and **running it under realistic conditions** (generalizability)



Procedure ---

- Experiment posted on **Mturk**
- Takes around **20 minutes**
- Participants paid **3 USD**
- Excluded participants who could not complete practice questions
- Total: **50-70** participants **out of 150**



Statistical Analysis: Linear Model ---


We use a **linear model** for each metric in each experiment:

- Response time
- Accuracy
- Satisfaction

Example – Model Size, Response Time:

- **Step 1:** Fit linear regression to predict response time from number of lines and number of output terms
- **Step 2:** Interpret coefficients as effects of number of lines and number of output terms on response time
 - ▶ The coefficient with the **largest magnitude** indicates which factor has the **strongest impact** on response time

🍷 Statistical Analysis: Multiple Hypothesis Testing ---



<https://xkcd.com/882/>


We use a Bonferroni correction

- Instead of $p < 0.05$, use: $p < (0.05 / \# \text{ comparisons})$

🍲 Results: Complexity Increases Response Time ---

Recipe Domain

V2 - Cognitive Chunks



Key Finding

Greater complexity results in longer response time for all kinds of complexity

🍲 Results: Type of Complexity Matters ---

		Response Time				
		Clinical		Recipe		
		Weight	P-Value	Weight	P-Value	
Model size	# Lines (V1)	01.17	<.0001	01.01	.0032	Significant in one domain
	# Terms (V1)	02.35	<.0001	01.57	.0378	
	Ver. (V1)	10.50	<.0001	04.11	.1210	
	Counter. (V1)	21.00	<.0001	13.70	<.0001	
Cognitive chunks	# Chunks (V2)	06.04	<.0001	05.88	<.0001	Significant in all domains
	Implicit (V2)	-13.00	<.0001	-07.93	0.0005	
	Ver. (V2)	16.30	<.0001	15.40	<.0001	
	Counter. (V2)	08.56	.0265	19.90	<.0001	
Variable repetitions	# Rep. (V3)	01.90	.2470	00.88	.4630	Significant in neither domain
	Ver. (V3)	13.00	.0035	13.70	<.0001	
	Counter. (V3)	20.30	<.0001	16.60	<.0001	

Key Finding

Response time for: cognitive chunks > model size > repeated terms



Results: Consistency - Domains, Tasks, Metrics ---

		Response Time			
		Clinical		Recipe	
		Weight	P-Value	Weight	P-Value
Model size	# Lines (V1)	01.17	<.0001	01.01	.0032
	# Terms (V1)	02.35	<.0001	01.57	.0378
	Ver. (V1)	10.50	<.0001	04.11	.1210
	Counter. (V1)	21.00	<.0001	13.70	<.0001
Cognitive chunks	# Chunks (V2)	06.04	<.0001	05.88	<.0001
	Implicit (V2)	-13.00	<.0001	-07.93	0.0005
	Ver. (V2)	16.30	<.0001	15.40	<.0001
	Counter. (V2)	08.56	.0265	19.90	<.0001
Variable repetitions	# Rep. (V3)	01.90	.2470	00.88	.4630
	Ver. (V3)	13.00	.0035	13.70	<.0001
	Counter. (V3)	20.30	<.0001	16.60	<.0001

For example:
Similar effect sizes,
both statistically
significant

Key Findings

Results consistent across domains, tasks and the response time and subjective difficulty metrics

🍲 Results: Counterfactuals Are Hard ---

The counterfactual task is much more challenging than simulation!

		Response Time			
		Clinical		Recipe	
		Weight	P-Value	Weight	P-Value
Model size	# Lines (V1)	01.17	<.0001	01.01	.0032
	# Terms (V1)	02.35	<.0001	01.57	.0378
	Ver. (V1)	10.50	<.0001	04.11	.1210
	Counter. (V1)	21.00	<.0001	13.70	<.0001
Cognitive chunks	# Chunks (V2)	06.04	<.0001	05.88	<.0001
	Implicit (V2)	-13.00	<.0001	-07.93	0.0005
	Ver. (V2)	16.30	<.0001	15.40	<.0001
	Counter. (V2)	08.56	.0265	19.90	<.0001
Variable repetitions	# Rep. (V3)	01.90	.2470	00.88	.4630
	Ver. (V3)	13.00	.0035	13.70	<.0001
	Counter. (V3)	20.30	<.0001	16.60	<.0001

In all experiments, longer response time than simulation

Key Finding

In all experiments, counterfactual tasks require longer response time than simulation tasks

Discussion: Paper 1 ---

- Consistent guidelines for interpretability
- Simplified tasks to measure interpretability
- Using Mturk workers as a proxy for domain experts

Key Takeaway

Different types of complexity affect **human-simulatability** differently, with cognitive chunks having the strongest effect

Paper 2 ---

🍷 Manipulating and Measuring Model Interpretability ---

Forough Poursabzi-Sangdeh, Daniel G. Goldstein, Jake M. Hofman

Jennifer Wortman Vaughan, Hanna Wallach

Microsoft Research



Motivation ---

- Interpretability as a latent property that can be manipulated or measured indirectly
- What are the factors through which it can be manipulated effectively?
- Bring HCI methods to interpretable ML since interpretability is defined by user experience

Research Questions:

- How well can people estimate what a model will predict?
- How much do people trust a model's predictions?
- How well can people detect when a model has made a sizable mistake?

Approach:

- Large-scale, pre-registered user studies to answer these questions in the context of linear regression models


🍷 Comparison to Paper 1 ---

- Studies **linear regression models** instead of decision sets
- Measures people's ability to make their **own predictions** in addition to forward simulation
- Uses **real-world housing dataset** and models optimized with data

5 Ways to Manipulate Interpretability ---



CLEAR-2: 2 features, transparent



BB-2: 2 features, black-box



Procedure ---

- Participants shown:
 - ▶ **Training:** 10 apartments
 - ▶ **Testing:** 12 apartments (this is the data they use)
- Participants paid 2.5 USD
- 750-1,250 participants per experiment

Each Trial:

Forward simulate model's prediction → View model's true prediction → Make own prediction

🍷 Statistical Analysis: Participant Specific Effects ---

A repeated measures experimental design

- Each participant makes many predictions
- Use a **mixed-effects model** to control for correlations between a participant's responses
- Assumes a random, participant-specific effect

Statistical Analysis: Multiple Hypothesis Testing ---

- **Pre-registering hypotheses** corresponds to deciding and publishing which analyses you will run **before collecting data**
- Reduces the probability that effects were discovered by chance

Example

<https://aspredicted.org/xy5s6.pdf>

"We will use 2-by-2 ANOVA for statistical analysis of the effect of number of features and model clarity on final deviation from model's prediction and simulation error."

Design Choices ---

- Randomized the order of the first 10 (normal) apartments and fixed the order of the last 2 (unusual)
- All participants are shown an identical set of apartments
- Each participant completed a single condition (between subjects design)

Fix sources of randomness

Can introduce bias

Randomize as much as possible

Increases variance

🍷 Results: Simulating Small, Transparent Models ---

Experiment 1: New York City prices




🍷 Results: No Difference in Trust or Prediction ---

Experiment 1: New York City prices



Trust (Deviation)



Prediction Error

Key Finding

🍷 Results: Clear Models Make Mistakes Worse ---

Experiment 1: New York City prices - Apartment 12 (1 bed, 3 bath)



Additional Experiments ---

Experiment 2: Scaled Prices

- Scaled down prices to better reflect national average
- **Same results** for simulation accuracy

Additional Experiments (cont.) ---

Experiment 3: Better Trust Metrics

- Scaled down prices to better reflect national average
- Better trust metrics
- **No significant difference in trust between models**

🍷 Additional Experiments (cont.) ---

Experiment 4: Attention Check

- Attention check for unusual features
- **People catch more errors** when explicitly asked to look for unusual inputs

Discussion: Paper 2 ---

Key findings from the experiments:

- Highlighting weird inputs helps catch errors
- Having people predict before seeing the model helped catch errors
- **Transparency actually makes people worse at catching errors**

Counterintuitive Result

More interpretable models (transparent, fewer features) led to worse error detection on unusual cases

General Discussion ---

Comparing Both Papers:

- **Paper 1:** Decision sets, controlled synthetic domains, cognitive complexity matters
- **Paper 2:** Linear models, real-world data, transparency can backfire

Key Lessons:

- Interpretability is multifaceted and context-dependent
- User studies are essential for evaluating interpretability
- Simplified models $\neq$ always better decision-making
- Different tasks require different evaluation metrics

Thank You!